

The analyst's first stop for any indicator.

A free, public, static threat-intelligence console for a T1/T2 SOC analyst. Paste an IP, hash, URL, domain, CVE or email and get one screen back: a verdict grounded in public reputation and vulnerability data, a ticket-ready escalation write-up, and one-click pivots to the services worth checking. Zero infrastructure, zero cost, zero accounts.

Positioning. The heavyweight open-source CTI platforms (OpenCTI, MISP, IntelOwl) need servers and databases. SOCDesk sits deliberately beneath them: the zero-infra, public-URL, analyst-speed layer an analyst actually opens mid-shift.

OPERATING PRINCIPLES — CONSTANT ACROSS EVERY LEG

- **\$0, free-tier only.** No paid infra, no paid APIs beyond free community tiers.
- **No-infrastructure.** Static site + Cloudflare Pages Functions. No database, no server to run.
- **Aggregator, not mirror.** Publish only clearly-redistributable data; reputation corpora are reached by user-clicked links, never held.
- **No secrets in the repo.** Collectors are keyless; enrichment keys live server-side in Cloudflare, never in git or the browser.
- **Privacy by default.** Analyst state (reviewed, watchlist, notable) is localStorage only. No accounts, no analytics, no tracking.
- **Strict CSP + defence in depth.** `default-src 'none'`, no inline script; every upstream string sanitised at collection and escaped again at render.
- **RADAR visual direction.** Dark, search-first analytics: an interactive real-data threat globe as hero, teal accent, data-table enrichment. Distinctive, not AI-slop. (Chart Room superseded.)

LEG 1 · CURRENT SPRINT

01 INDICATOR LOOKUP & ESCALATION

CORE · LIVE

The 99% workflow: paste an indicator → live reputation verdict → screenshot-grade evidence card → paste into the escalation email. This Leg is the product.

INPUT & DETECTION

- ✓ **Omnibox** with type auto-detect + refang
(evil[.]com, hxxp)
- ✓ **Bulk lookup** — many indicators, one table, CSV/JSON/defanged export
- ✓ **Bookmarklet** — fire a lookup from any page
- ✓ **Shareable #q= deep links** — "look at this" is a URL

LIVE ENRICHMENT · /API/ENRICH

- ✓ **IP** → AbuseIPDB · VirusTotal · GreyNoise · ipinfo geo/ASN
- ✓ **Hash** → VirusTotal · MalwareBazaar
- ✓ **URL / domain** → VirusTotal · urlscan verdict + screenshot
- ✓ **Consensus tally** — "N of M sources flagged"; SOCDesk counts, never asserts a verdict

OUTPUT — THE DIFFERENTIATOR

- ✓ **Escalation card** — the analyst copies it into the ticket; no SOCDesk recommendation (their call), button reads COPY
- ✓ **Type-aware pivots**, trimmed to the T1 workflow set (no fluff vendors)
- ✓ **CVE** — authoritative, attributed: CISA-confirmed exploited (KEV) + CVSS + EPSS
- ✓ **Honest states** — "0 of M flagged" is never dressed as a clearance

REACH — THE RECORDED-FUTURE MOVE

- ✓ **Browser extension v1** — right-click lookup + toolbar popup over the live API (built, permission-reviewed)
- **v2 · inline page annotation** — highlight indicators on any page (broad perms — later)
- **Bookmarklet selection-capture** — no-install fallback

LIVE at socdesk.io. In progress: the RADAR visual rebuild (React + shadcn), shipping the extension (Unlisted), and analyst dogfood. Verdict output is the consensus tally — see docs/VERDICT-LANGUAGE.md.

02 THREAT INTELLIGENCE FEED

CORE · EVOLVE, DON'T REBUILD

Situational awareness for the shift: what's happening, what's exploited, what's relevant to us. The pipeline is already strong and tested — Leg 2 is additive, not a rewrite.

ALREADY BUILT (THE FOUNDATION)

- ✓ **Collector pipeline** — KEV, NVD, EPSS, Ransomware.live, RSS pool; fault-isolated, schema-gated with last-known-good
- ✓ **CVE triage** — full KEV catalogue joined with CVSS + EPSS, risk-sorted
- ✓ **Feed work-queue** — explainable relevance scoring, claim grouping
- ✓ **ATT&CK actor/malware profiles + RELATED** entity index
- ✓ **Collection health**, source registry, trends

PLANNED EVOLUTION (ADDITIVE)

- **AI Daily Brief** — Framework local LLM writes the shift brief (passthrough already wired)
- **Client / sector relevance** — "affects our clients" filter, the MSSP differentiator
- **Wave-2 collectors** — FeodoTracker, C2IntelFeeds, PhishTank, APTnotes
- **Push / alerting** — shift digest to Slack/email
- **Decouple data-refresh from site-deploy** for higher cadence

DESIGN DECISIONS TO LOCK FIRST

- ▶ What the feed is *for* (brief vs. relevance-filter vs. hunt-feed)
- ▶ Data backbone (static vs. KV/R2) · watchlist location · alerting scope

Approach: a proper design pass before building, on top of the existing pipeline. The early "garbage" was the UI/identity — long since fixed — not the data layer.

SOCDesk stops being only an aggregator and starts producing first-party intelligence: our own honeypot sensors observing real attacks, published as a live feed that loops back into Legs 1 & 2.

CORE FUNCTIONS

- Own honeypot sensor(s) — Cowrie-class, capturing live SSH/Telnet/scan activity
- Live attack telemetry & globe visualisation (mockup: `design/mockups/h-sensor.html`)
- First-party IOC feed — observed attacker infrastructure, feeding Leg-1 verdicts and Leg-2 feed
- Trend surfacing — emerging scan campaigns, credential-spray patterns

SECURITY NON-NEGOTIABLES

- No Tailscale on the sensor — no path from the honeypot into the tailnet
- Cloud-plane egress deny by default
- Publish /24s, not full IPs — no doxxing of compromised hosts
- Allowlist published tokens; rebuild from cloud-init, never snapshots

GATE

- Resolve the employer-IP-assignment question first — it gates this Leg (and the repo-private decision)

SEQUENCING & CURRENT STATUS

- ✓ Done: Leg 1 is LIVE at socdesk.io — live multi-source enrichment, the consensus tally, and the browser extension v1.
- Now: the RADAR visual rebuild (React + shadcn) — final branding, a hero-scale interactive globe, a designed escalation card; ship the extension (Unlisted); analyst dogfood.
- Then: settle the employer-IP question → repo public vs. private; Leg 2 threat-intel feed on the existing pipeline.
- Later: Leg 3 sensor, once the IP gate clears.